

Register Number: 192425317

Name: M. Sahasra

CO3 – AT1

MLA0201 – Fundamentals of Machine Learning

Experiments

1) Customer Segmentation using K-Means and PCA: Load the Mall Customers dataset and preprocess the data. Apply the K-Means algorithm to identify customer clusters, then use PCA to reduce the data to two dimensions and visualize the clusters. Analyze the optimal number of clusters using the Elbow Method and Silhouette Score.

Answer)

Definitions:

Unsupervised Learning:

Unsupervised learning is a machine learning technique in which the algorithm learns patterns from unlabelled data. There is no predefined target/output variable.

Clustering:

Clustering is an unsupervised learning technique that divides data into groups called clusters, where similar data points are placed together.

K-Means Clustering:

K-Means is a clustering algorithm that divides data into K clusters.

The basic steps are:

1. Select the number of clusters, K.
2. Initialize K centroids.

3. Assign each data point to its nearest centroid.
4. Recalculate the centroids.
5. Repeat until the clusters stabilize.

Centroid:

A centroid is the central point representing a cluster. In K-Means, the centroid is calculated as the mean of all points belonging to that cluster.

PCA — Principal Component Analysis:

PCA is a dimensionality reduction technique that transforms a large number of correlated features into a smaller number of principal components while preserving as much information as possible.

Elbow Method:

The Elbow Method helps determine a suitable value of K by plotting Within-Cluster Sum of Squares (WCSS) against different values of K.

The point where the decrease in WCSS becomes slower is called the elbow point.

Silhouette Score:

Silhouette Score measures how well each data point fits within its assigned cluster.

Its value ranges from:

-1 to +1

A value closer to +1 generally indicates better-defined clusters.

Objective:

To segment mall customers into different groups using K-Means clustering, determine the optimal number of clusters using the Elbow Method and Silhouette Score, and visualize the clusters using PCA.

Python Code:

```
1 # =====
2 # EXPERIMENT 1
3 # Customer Segmentation using K-Means and PCA
4 # =====
5
6 # Import libraries
7 import pandas as pd
8 import numpy as np
9 import matplotlib.pyplot as plt
10
11 from sklearn.cluster import KMeans
12 from sklearn.preprocessing import StandardScaler
13 from sklearn.decomposition import PCA
14 from sklearn.metrics import silhouette_score
15
16 # -----
17 # STEP 1: Load Mall Customers Dataset
18 # -----
19
20 url = "https://raw.githubusercontent.com/sharnaroshan/Clustering-of-Mall-Customers/master/Mall_Customers.csv"
21
22 df = pd.read_csv(url)
23
24 print("First 5 rows of the dataset:")
25 print(df.head())
26
27 print("\nDataset Shape:")
28 print(df.shape)
29
30 # -----
31 # STEP 2: Data Preprocessing
32 # -----
33
34 # Check missing values
35 print("\nMissing Values:")
36 print(df.isnull().sum())
37
38 # Select numerical features
39 features = [
40     "Age",
41     "Annual Income (k$)",
42     "Spending Score (1-100)"
43 ]
44
45 X = df[features]
46
47 # Standardize the data
48 scaler = StandardScaler()
49 X_scaled = scaler.fit_transform(X)
50
51 print("\nStandardized data (first 5 rows):")
52 print(X_scaled[:5])
53
54 # -----
55 # STEP 3: Elbow Method
56 # -----
57
58 wcss = []
59
60 for k in range(2, 11):
61     kmeans = KMeans(
62         n_clusters=k,
63         random_state=42,
64         n_init=10
65     )
66
67     kmeans.fit(X_scaled)
68     wcss.append(kmeans.inertia_)
69
70 # Plot Elbow Curve
71 plt.figure(figsize=(8, 5))
72 plt.plot(
73     range(2, 11),
74     wcss,
75     marker='o'
76 )
77
78 plt.title("Elbow Method for Optimal K")
79 plt.xlabel("Number of Clusters (K)")
80 plt.ylabel("WCSS")
81 plt.grid(True)
82 plt.show()
```

```

83
84 # -----
85 # STEP 4: Silhouette Score
86 # -----
87
88 silhouette_scores = []
89
90 for k in range(2, 11):
91
92     kmeans = KMeans(
93         n_clusters=k,
94         random_state=42,
95         n_init=10
96     )
97
98     labels = kmeans.fit_predict(X_scaled)
99
100     score = silhouette_score(X_scaled, labels)
101
102     silhouette_scores.append(score)
103
104 # Display scores
105 print("\nSilhouette Scores:")
106
107 for k, score in zip(range(2, 11), silhouette_scores):
108     print(f"K = {k}, Silhouette Score = {score:.4f}")
109
110 # Find best K according to silhouette score
111 best_k = range(2, 11)[np.argmax(silhouette_scores)]
112
113 print("\nBest K according to Silhouette Score:", best_k)
114
115 # Plot silhouette scores
116 plt.figure(figsize=(8, 5))
117
118 plt.plot(
119     range(2, 11),
120     silhouette_scores,
121     marker='o'
122 )
123
124 plt.title("Silhouette Score for Different K Values")
125 plt.xlabel("Number of Clusters (K)")
126 plt.ylabel("Silhouette Score")
127 plt.grid(True)
128 plt.show()
129
130 # -----
131 # STEP 5: Apply K-Means using Best K
132 # -----
133
134 kmeans = KMeans(
135     n_clusters=best_k,
136     random_state=42,
137     n_init=10
138 )
139
140 clusters = kmeans.fit_predict(X_scaled)
141
142 # Add cluster labels to original dataframe
143 df["Cluster"] = clusters
144
145 print("\nCustomer Cluster Assignment:")
146 print(df.head(10))
147
148 # -----
149 # STEP 6: Apply PCA
150 # -----
151
152 pca = PCA(n_components=2)
153
154 X_pca = pca.fit_transform(X_scaled)
155
156 print("\nPCA Explained Variance Ratio:")
157 print(pca.explained_variance_ratio_)
158
159 print(
160     "\nTotal Variance Explained:",
161     round(sum(pca.explained_variance_ratio_) * 100, 2),
162     "%"
163 )
164

```

```

165 # -----
166 # STEP 7: Visualize Clusters using PCA
167 # -----
168
169 plt.figure(figsize=(9, 6))
170
171 scatter = plt.scatter(
172     X_pca[:, 0],
173     X_pca[:, 1],
174     c=clusters,
175     cmap="viridis",
176     s=50
177 )
178
179 plt.title("Customer Segmentation using K-Means + PCA")
180 plt.xlabel("Principal Component 1")
181 plt.ylabel("Principal Component 2")
182
183 plt.colorbar(scatter, label="Cluster")
184
185 plt.grid(True)
186 plt.show()
187
188 # -----
189 # STEP 8: Cluster Summary
190 # -----
191
192 print("\nCluster Summary:")
193
194 summary = df.groupby("Cluster")[features].mean()
195
196 print(summary)
197
198 # -----
199 # FINAL RESULT
200 # -----
201
202 print("\n=====")
203 print("EXPERIMENT 1 COMPLETED")
204 print("=====")
205 print("Optimal K:", best_k)
206 print("Number of Customers:", len(df))
207 print(
208     "PCA Variance Explained:",
209     round(sum(pca.explained_variance_ratio_) * 100, 2),
210     "%"
211 )

```

Output:

First 5 rows of the dataset:

CustomerID	Genre	Age	Annual Income (k\$)	Spending Score (1-100)
0	1 Male	19	15	39
1	2 Male	21	15	81
2	3 Female	20	16	6
3	4 Female	23	16	77
4	5 Female	31	17	40

Dataset Shape:
(200, 5)

Missing Values:

CustomerID	0
Genre	0
Age	0
Annual Income (k\$)	0
Spending Score (1-100)	0


dtype: int64

Standardized data (first 5 rows):

```

[[-1.42456879 -1.73899919 -0.43480148]
 [-1.28103541 -1.73899919  1.19570407]
 [-1.3528021  -1.70082976 -1.71591298]
 [-1.13750203 -1.70082976  1.04041783]
 [-0.56336851 -1.66266033 -0.39597992]]


```



Silhouette Scores:

...
K = 2, Silhouette Score = 0.3355
K = 3, Silhouette Score = 0.3578
K = 4, Silhouette Score = 0.4040
K = 5, Silhouette Score = 0.4166
K = 6, Silhouette Score = 0.4284
K = 7, Silhouette Score = 0.4172
K = 8, Silhouette Score = 0.4082
K = 9, Silhouette Score = 0.4177
K = 10, Silhouette Score = 0.4066

Best K according to Silhouette Score: 6



Customer Cluster Assignment:

	CustomerID	Genre	Age	Annual Income (k\$)	Spending Score (1-100)	\
0	1	Male	19	15	39	
1	2	Male	21	15	81	
2	3	Female	20	16	6	
3	4	Female	23	16	77	
4	5	Female	31	17	40	
5	6	Female	22	17	76	
6	7	Female	35	18	6	
7	8	Female	23	18	94	
8	9	Male	64	19	3	
9	10	Female	30	19	72	


```

Cluster
0      4
... 1      4
2      5
3      4
4      5
5      4
6      5
7      4
8      5
9      4

```

PCA Explained Variance Ratio:
[0.44266167 0.33308378]

Total Variance Explained: 77.57 %



Cluster Summary:

	Age	Annual Income (k\$)	Spending Score (1-100)
Cluster			
0	56.333333	54.266667	49.066667
1	26.794872	57.102564	48.128205
2	41.939394	88.939394	16.969697
3	32.692308	86.538462	82.128205
4	25.000000	25.260870	77.608696
5	45.523810	26.285714	19.380952

=====

EXPERIMENT 1 COMPLETED

=====

Optimal K: 6

Number of Customers: 200

PCA Variance Explained: 77.57 %

Result:

Thus, K-Means clustering successfully segmented the mall customers into groups, and PCA was used to represent the clusters in two dimensions. The Elbow Method and Silhouette Score were used to determine an appropriate number of clusters.

2) Clustering using EM Algorithm and Gaussian Mixture Model: Load the Digits dataset from Scikit-learn and preprocess the features. Implement the Gaussian Mixture Model (EM Algorithm) to cluster the data and compare its performance with K-Means using appropriate clustering evaluation metrics. Visualize the clusters after applying PCA.

Answer)**Definitions:****Gaussian Mixture Model — GMM:**

A Gaussian Mixture Model is a probabilistic clustering model that assumes the data is generated from a mixture of several Gaussian probability distributions.

Each Gaussian distribution represents a cluster.

EM Algorithm:

EM stands for Expectation-Maximization.

It is an iterative optimization algorithm used to estimate parameters when some information is hidden or unknown.

It has two main steps:

E-Step — Expectation:

Estimate the probability that each data point belongs to each cluster.

M-Step — Maximization:

Update the model parameters based on those probabilities.

These steps are repeated until the model converges.

Hard Clustering:

In hard clustering, each data point belongs to exactly one cluster.

K-Means is an example of hard clustering.

Soft Clustering:

In soft clustering, a data point can belong to multiple clusters with different probabilities.

GMM performs soft clustering.

K-Means vs GMM:

K-Means	GMM
Distance-based	Probability-based
Hard clustering	Soft clustering
Uses centroids	Uses Gaussian distributions
Assigns one cluster	Provides cluster probabilities
Simpler	More flexible

Objective:

To cluster handwritten digit data using Gaussian Mixture Model with the EM algorithm, compare its performance with K-Means, and visualize the resulting clusters using PCA.

Python Code:

```
1 # =====
2 # EXPERIMENT 2
3 # Clustering using EM Algorithm and Gaussian Mixture Model
4 # =====
5
6 # Import libraries
7 import numpy as np
8 import pandas as pd
9 import matplotlib.pyplot as plt
10
11 from sklearn.datasets import load_digits
12 from sklearn.preprocessing import StandardScaler
13 from sklearn.cluster import KMeans
14 from sklearn.mixture import GaussianMixture
15 from sklearn.decomposition import PCA
16 from sklearn.metrics import (
17     silhouette_score,
18     adjusted_rand_score,
19     normalized_mutual_info_score
20 )
21
22 # -----
23 # STEP 1: Load Digits Dataset
24 # -----
25
26 digits = load_digits()
27
28 X = digits.data
29 y = digits.target
30
31 print("Dataset Shape:")
32 print(X.shape)
33
34 print("\nNumber of Classes:")
35 print(len(np.unique(y)))
36
37 print("\nFirst 5 target values:")
38 print(y[:5])
39
40 # -----
41 # STEP 2: Standardize Features
42 # -----
43
44 scaler = StandardScaler()
45
46 X_scaled = scaler.fit_transform(X)
47
48 print("\nStandardized data shape:")
49 print(X_scaled.shape)
50
51 # -----
52 # STEP 3: Apply K-Means
53 # -----
54
55 # Digits dataset contains 10 digit classes
56 n_clusters = 10
57
58 kmeans = KMeans(
59     n_clusters=n_clusters,
60     random_state=42,
61     n_init=10
62 )
63
64 kmeans_labels = kmeans.fit_predict(X_scaled)
65
66 # -----
67 # STEP 4: Apply Gaussian Mixture Model
68 # -----
69
70 gmm = GaussianMixture(
71     n_components=n_clusters,
72     covariance_type="full",
73     random_state=42
74 )
75
76 gmm_labels = gmm.fit_predict(X_scaled)
77
78 # -----
79 # STEP 5: Evaluate K-Means
80 # -----
81
82 kmeans_silhouette = silhouette_score(
```

```

83     X_scaled,
84     kmeans_labels
85 )
86
87 kmeans_ari = adjusted_rand_score(
88     y,
89     kmeans_labels
90 )
91
92 kmeans_nmi = normalized_mutual_info_score(
93     y,
94     kmeans_labels
95 )
96
97 # -----
98 # STEP 6: Evaluate GMM
99 # -----
100
101 gmm_silhouette = silhouette_score(
102     X_scaled,
103     gmm_labels
104 )
105
106 gmm_ari = adjusted_rand_score(
107     y,
108     gmm_labels
109 )
110
111 gmm_nmi = normalized_mutual_info_score(
112     y,
113     gmm_labels
114 )
115
116 # -----
117 # STEP 7: Display Comparison
118 # -----
119
120 results = pd.DataFrame({
121     "Algorithm": [
122         "K-Means",
123         "Gaussian Mixture Model"
124     ],
125     "Silhouette Score": [
126         kmeans_silhouette,
127         gmm_silhouette
128     ],
129     "Adjusted Rand Index": [
130         kmeans_ari,
131         gmm_ari
132     ],
133     "Normalized Mutual Information": [
134         kmeans_nmi,
135         gmm_nmi
136     ]
137 })
138
139 print("\n=====")
140 print("CLUSTERING PERFORMANCE COMPARISON")
141 print("=====")
142
143 print(results.to_string(index=False))
144
145 # -----
146 # STEP 8: PCA for Visualization
147 # -----
148
149 pca = PCA(n_components=2)
150
151 X_pca = pca.fit_transform(X_scaled)
152
153 print("\nPCA Explained Variance Ratio:")
154 print(pca.explained_variance_ratio_)
155
156 print(
157     "Total Variance Explained:",
158     round(sum(pca.explained_variance_ratio_) * 100, 2),
159     "%"
160 )

```

```

164
165 # -----
166 # STEP 9: Visualize K-Means Clusters
167 # -----
168
169 plt.figure(figsize=(9, 6))
170
171 scatter = plt.scatter(
172     X_pca[:, 0],
173     X_pca[:, 1],
174     c=kmeans_labels,
175     cmap="tab10",
176     s=15
177 )
178
179 plt.title("Digits Clustering using K-Means + PCA")
180 plt.xlabel("Principal Component 1")
181 plt.ylabel("Principal Component 2")
182 plt.colorbar(scatter, label="Cluster")
183 plt.grid(True)
184 plt.show()
185
186 # -----
187 # STEP 10: Visualize GMM Clusters
188 # -----
189
190 plt.figure(figsize=(9, 6))
191
192 scatter = plt.scatter(
193     X_pca[:, 0],
194     X_pca[:, 1],
195     c=gmm_labels,
196     cmap="tab10",
197     s=15
198 )
199
200 plt.title("Digits Clustering using GMM + PCA")
201 plt.xlabel("Principal Component 1")
202 plt.ylabel("Principal Component 2")
203 plt.colorbar(scatter, label="Cluster")
204 plt.grid(True)
205 plt.show()
206
207 # -----
208 # STEP 11: Display GMM Cluster Probabilities
209 # -----
210
211 probabilities = gmm.predict_proba(X_scaled)
212
213 print("\nGMM Cluster Probabilities for First 5 Samples:")
214
215 print(
216     pd.DataFrame(
217         probabilities[:5],
218         columns=[f"Cluster_{i}" for i in range(n_clusters)]
219     ).round(3)
220 )
221
222 # -----
223 # FINAL RESULT
224 # -----
225
226 print("\n=====")
227 print("EXPERIMENT 2 COMPLETED")
228 print("=====")
229
230 print("K-Means Silhouette Score:",
231       round(kmeans_silhouette, 4))
232
233 print("GMM Silhouette Score:",
234       round(gmm_silhouette, 4))
235
236 print("K-Means ARI:",
237       round(kmeans_ari, 4))
238
239 print("GMM ARI:",
240       round(gmm_ari, 4))

```

Output:

```
Dataset Shape:
... (1797, 64)

Number of Classes:
10


First 5 target values:
[0 1 2 3 4]

Standardized data shape:
(1797, 64)


=====
CLUSTERING PERFORMANCE COMPARISON
=====
```

Algorithm	Silhouette Score	Adjusted Rand Index	Normalized Mutual Information
K-Means	0.139377	0.534407	0.671244
Gaussian Mixture Model	0.117905	0.546745	0.690155

```
PCA Explained Variance Ratio:
[0.12033916 0.09561054]
Total Variance Explained: 21.59 %
```



...



GMM Cluster Probabilities for First 5 Samples:

	Cluster_0	Cluster_1	Cluster_2	Cluster_3	Cluster_4	Cluster_5 \
0	0.0	0.0	0.0	0.0	0.0	1.0
1	0.0	0.0	0.0	0.0	0.0	0.0
2	0.0	0.0	0.0	0.0	0.0	0.0
3	0.0	0.0	0.0	0.0	0.0	0.0
4	0.0	0.0	1.0	0.0	0.0	0.0

	Cluster_6	Cluster_7	Cluster_8	Cluster_9
0	0.0	0.0	0.0	0.0
1	0.0	0.0	0.0	1.0
2	0.0	0.0	0.0	1.0
3	0.0	1.0	0.0	0.0
4	0.0	0.0	0.0	0.0

```
=====
EXPERIMENT 2 COMPLETED
=====
K-Means Silhouette Score: 0.1394
GMM Silhouette Score: 0.1179
K-Means ARI: 0.5344
GMM ARI: 0.5467
```

Result:

Thus, the Gaussian Mixture Model using the EM algorithm was successfully applied to the Digits dataset and its clustering performance was compared with K-Means using clustering evaluation metrics. PCA was used to visualize the clusters in two dimensions.

3) Dimensionality Reduction using PCA, Factor Analysis, and ICA: Load the Wine dataset from Scikit-learn and standardize the features. Apply Principal Component Analysis (PCA), Factor Analysis (FA), and Independent Component Analysis (ICA) to reduce the dimensionality and compare their outputs. Discuss how these techniques help address the Curse of Dimensionality through visualizations and component analysis.

Answer)

Definitions:

Dimensionality Reduction:

Dimensionality reduction is the process of reducing the number of features in a dataset while retaining the most important information.

For example:

13 features $\rightarrow$ 2 components

This makes data easier to visualize and can reduce computational complexity.

Curse of Dimensionality:

The Curse of Dimensionality refers to problems that occur when the number of features becomes very large.

High-dimensional data can lead to:

- Increased computational cost
- More complex models
- Difficulty in visualization
- Increased noise
- Poor performance in some machine learning algorithms

PCA:

Principal Component Analysis transforms the original features into a set of uncorrelated principal components.

The first component captures the maximum possible variance.

Factor Analysis:

Factor Analysis attempts to explain observed variables using a smaller number of hidden or latent factors.

It is commonly used when the observed variables are believed to be influenced by underlying factors.

ICA:

Independent Component Analysis transforms the original data into components that are as statistically independent as possible.

ICA is useful when the observed data is assumed to be a mixture of independent underlying sources.

PCA vs FA vs ICA:

Technique	Main Idea
PCA	Maximum variance
Factor Analysis	Find latent factors
ICA	Find statistically independent components

Objective:

To reduce the dimensionality of the Wine dataset using PCA, Factor Analysis and ICA, compare the resulting components, visualize the reduced data, and understand how dimensionality reduction helps address the Curse of Dimensionality.

Python Code:

```
1 # =====
2 # EXPERIMENT 3
3 # Dimensionality Reduction using PCA, Factor Analysis and ICA
4 # =====
5
6 # Import libraries
7 import numpy as np
8 import pandas as pd
9 import matplotlib.pyplot as plt
10
11 from sklearn.datasets import load_wine
12 from sklearn.preprocessing import StandardScaler
13 from sklearn.decomposition import (
14     PCA,
15     FactorAnalysis,
16     FastICA
17 )
18
19 # -----
20 # STEP 1: Load Wine Dataset
21 # -----
22
23 wine = load_wine()
24
25 X = wine.data
26 y = wine.target
27
28 feature_names = wine.feature_names
29
30 print("Dataset Shape:")
31 print(X.shape)
32
33 print("\nFeature Names:")
34 print(feature_names)
35
36 print("\nFirst 5 Target Values:")
37 print(y[:5])
38
39 # -----
40 # STEP 2: Create DataFrame
41 # -----
42
43 df = pd.DataFrame(
44     X,
45     columns=feature_names
46 )
47
48 print("\nFirst 5 rows:")
49 print(df.head())
50
51 # -----
52 # STEP 3: Standardize Features
53 # -----
54
55 scaler = StandardScaler()
56
57 X_scaled = scaler.fit_transform(X)
58
59 print("\nStandardized data:")
60 print(X_scaled[:5])
61
62 # -----
63 # STEP 4: PCA
64 # -----
65
66 pca = PCA(n_components=2)
67
68 X_pca = pca.fit_transform(X_scaled)
69
70 print("\n=====")
71 print("PCA RESULTS")
72 print("=====")
73
74 print("PCA Component Shape:")
75 print(X_pca.shape)
76
77 print("\nPCA Explained Variance Ratio:")
78 print(pca.explained_variance_ratio_)
79
80 print(
81     "\nTotal Variance Explained by PCA:",
82     round(sum(pca.explained_variance_ratio_) * 100, 2),
83     ..
84 )
```

```

83     "%"
84 )
85
86 # -----
87 # STEP 5: Factor Analysis
88 # -----
89
90 fa = FactorAnalysis(
91     n_components=2,
92     random_state=42
93 )
94
95 X_fa = fa.fit_transform(X_scaled)
96
97 print("\n=====")
98 print("FACTOR ANALYSIS RESULTS")
99 print("=====")
100
101 print("Factor Analysis Component Shape:")
102 print(X_fa.shape)
103
104 # Calculate approximate variance representation
105 fa_variance = np.var(X_fa, axis=0)
106
107 print("\nVariance of Factor Components:")
108 print(fa_variance)
109
110 # -----
111 # STEP 6: Independent Component Analysis
112 # -----
113
114 ica = FastICA(
115     n_components=2,
116     random_state=42,
117     max_iter=2000,
118     whiten="unit-variance"
119 )
120
121 X_ica = ica.fit_transform(X_scaled)
122
123 print("\n=====")
124 print("ICA RESULTS")
125 print("=====")
126
127 print("ICA Component Shape:")
128 print(X_ica.shape)
129
130 print("\nICA Component Variance:")
131 print(np.var(X_ica, axis=0))
132
133 # -----
134 # STEP 7: Create Comparison DataFrame
135 # -----
136
137 comparison = pd.DataFrame({
138     "Method": [
139         "PCA",
140         "Factor Analysis",
141         "ICA"
142     ],
143     "Component 1 Variance": [
144         np.var(X_pca[:, 0]),
145         np.var(X_fa[:, 0]),
146         np.var(X_ica[:, 0])
147     ],
148     "Component 2 Variance": [
149         np.var(X_pca[:, 1]),
150         np.var(X_fa[:, 1]),
151         np.var(X_ica[:, 1])
152     ]
153 })
154
155 print("\n=====")
156 print("DIMENSIONALITY REDUCTION COMPARISON")
157 print("=====")
158
159 print(comparison.to_string(index=False))
160
161 # -----

```

```

164 # STEP 8: PCA Visualization
165 # -----
166
167 plt.figure(figsize=(8, 6))
168
169 scatter = plt.scatter(
170     X_pca[:, 0],
171     X_pca[:, 1],
172     c=y,
173     cmap="viridis",
174     s=45
175 )
176
177 plt.title("Wine Dataset - PCA")
178 plt.xlabel("Principal Component 1")
179 plt.ylabel("Principal Component 2")
180 plt.colorbar(scatter, label="Wine Class")
181 plt.grid(True)
182 plt.show()
183
184 # -----
185 # STEP 9: Factor Analysis Visualization
186 # -----
187
188 plt.figure(figsize=(8, 6))
189
190 scatter = plt.scatter(
191     X_fa[:, 0],
192     X_fa[:, 1],
193     c=y,
194     cmap="viridis",
195     s=45
196 )
197
198 plt.title("Wine Dataset - Factor Analysis")
199 plt.xlabel("Factor 1")
200 plt.ylabel("Factor 2")
201 plt.colorbar(scatter, label="Wine Class")
202 plt.grid(True)
203 plt.show()
204
205 # -----
206 # STEP 10: ICA Visualization
207 # -----
208
209 plt.figure(figsize=(8, 6))
210
211 scatter = plt.scatter(
212     X_ica[:, 0],
213     X_ica[:, 1],
214     c=y,
215     cmap="viridis",
216     s=45
217 )
218
219 plt.title("Wine Dataset - Independent Component Analysis")
220 plt.xlabel("Independent Component 1")
221 plt.ylabel("Independent Component 2")
222 plt.colorbar(scatter, label="Wine Class")
223 plt.grid(True)
224 plt.show()
225
226 # -----
227 # STEP 11: Compare Original and Reduced Dimensions
228 # -----
229
230 print("\n=====")
231 print("DIMENSION REDUCTION")
232 print("=====")
233
234 print("Original number of features:", X.shape[1])
235
236 print("PCA features:", X_pca.shape[1])
237
238 print("Factor Analysis features:", X_fa.shape[1])
239
240 print("ICA features:", X_ica.shape[1])
241
242 print("\nDimensionality reduction:")
243
244 print(
245     f"PCA: {X.shape[1]} -> {X_pca.shape[1]}"

```

```

246 )
247
248 print(
249     f"Factor Analysis: {X.shape[1]} -> {X_fa.shape[1]}"
250 )
251
252 print(
253     f"ICA: {X.shape[1]} -> {X_ica.shape[1]}"
254 )
255
256 # -----
257 # FINAL RESULT
258 # -----
259
260 print("\n=====")
261 print("EXPERIMENT 3 COMPLETED")
262 print("=====")
263
264 print(
265     "PCA retained",
266     round(sum(pca.explained_variance_ratio_) * 100, 2),
267     "% of the variance using 2 components."
268 )
269
270 print(
271     "All three techniques successfully reduced",
272     X.shape[1],
273     "features to 2 components."
274 )

```

Output:

```

Dataset Shape:
(178, 13)

...
Feature Names:
['alcohol', 'malic_acid', 'ash', 'alcalinity_of_ash', 'magnesium', 'total_phenols', 'flavanoids', 'nonflavanoid_phenols', 'proanthocyanins', 'color_intensity', 'hue', 'od280/od315_of_diluted_wines', 'proline']

First 5 Target Values:
[0 0 0 0 0]

First 5 rows:
alcohol malic_acid ash alcalinity_of_ash magnesium total_phenols \
0 14.23 1.71 2.43 15.6 127.0 2.80
1 13.20 1.78 2.14 11.2 100.0 2.65
2 13.16 2.36 2.67 18.6 101.0 2.80
3 14.37 1.95 2.50 16.8 113.0 3.85
4 13.24 2.59 2.87 21.0 118.0 2.80

flavanoids nonflavanoid_phenols proanthocyanins color_intensity hue \
0 3.06 0.28 2.29 5.64 1.04
1 2.76 0.26 1.28 4.38 1.05
2 3.24 0.30 2.81 5.68 1.03
3 3.49 0.24 2.18 7.80 0.86
4 2.69 0.39 1.82 4.32 1.04

od280/od315_of_diluted_wines proline
0 3.92 1065.0
1 3.40 1050.0
2 3.17 1105.0
3 3.45 1400.0
4 2.93 735.0

Standardized data:
[[ 1.51861254 -0.5622498 0.23205254 -1.16959318 1.91390522 0.80899739
 1.03481896 -0.65956311 1.22488398 0.25171685 0.36217728 1.84791957
 1.01300893]
 [ 0.24628963 -0.49941338 -0.82799632 -2.49084714 0.01814502 0.56864766
 0.73362894 -0.82071924 -0.54472099 -0.29332133 0.40605066 1.1134493
 0.96524152]
 [ 0.19667903 0.02123125 1.10933436 -0.2687382 0.08835836 0.80899739
 1.21553297 -0.49840699 2.13596773 0.26901965 0.31830389 0.78858745
 1.39514818]
 [ 1.69154964 -0.34681064 0.4879264 -0.80925118 0.93091845 2.49144552
 1.46652465 -0.98187536 1.03215473 1.18606801 -0.42754369 1.18407144
 2.33457383]
 [ 0.29570023 0.22769377 1.84040254 0.45194578 1.28198515 0.80899739
 0.66335127 0.22679555 0.40140444 -0.31927553 0.36217728 0.44960118
 -0.03787401]]

=====
PCA RESULTS
=====
PCA Component Shape:
(178, 2)

PCA Explained Variance Ratio:
[0.36198848 0.1920749 ]

Total Variance Explained by PCA: 55.41 %

```

```
***
=====
FACTOR ANALYSIS RESULTS
=====
Factor Analysis Component Shape:
(178, 2)
```

```
Variance of Factor Components:
[0.95605362 0.87119034]
```



ICA RESULTS


```
ICA Component Shape:
(178, 2)
```

```
ICA Component Variance:
[1. 1.]
```

DIMENSIONALITY REDUCTION COMPARISON

Method	Component 1	Variance	Component 2	Variance
PCA		4.705850		2.496974
Factor Analysis		0.956054		0.871190
ICA		1.000000		1.000000





```
=====
DIMENSION REDUCTION
=====
Original number of features: 13
PCA features: 2
Factor Analysis features: 2
ICA features: 2

Dimensionality reduction:
PCA: 13 -> 2
Factor Analysis: 13 -> 2
ICA: 13 -> 2

=====
EXPERIMENT 3 COMPLETED
=====
PCA retained 55.41 % of the variance using 2 components.
All three techniques successfully reduced 13 features to 2 components.
```

Result:

Thus, the Wine dataset was successfully reduced from 13 dimensions to 2 dimensions using PCA, Factor Analysis and ICA. PCA focuses on preserving maximum variance, Factor Analysis identifies latent factors, and ICA attempts to identify statistically independent components. These dimensionality reduction techniques help make high-dimensional data easier to visualize and analyze and help address issues associated with the Curse of Dimensionality.